

Lab Assignment No.: 7 Feature Engineering

Apply Preprocessing steps/ Feature Engineering like handling missing values, transformations etc on the given dataset. Write your observation in conclusion

Theory

All machine learning algorithms use some input data to generate outputs. Input data contains many features which may not be in proper form to be given to the model directly. It needs some kind of processing and here feature engineering helps. Feature engineering fulfils mainly two goals:

It prepares the input dataset in the form which is required for a specific model or machine learning algorithm. Feature engineering helps in improving the performance of machine learning models magically.

The main feature engineering techniques are:

1. Missing data imputation
2. Categorical encoding
3. Variable transformation
4. Outlier engineering
5. Date and time engineering

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
df = pd.read_csv("adult.csv")
```

```
df.head()
```

	age	workclass	fnlwgt	education	educational-num	marital-status \
0	25	Private	226802	11th	7	Never-married
1	38	Private	89814	HS-grad	9	Married-civ-spouse
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse
3	44	Private	160323	Some-college	10	Married-civ-spouse
4	18	?	103497	Some-college	10	Never-

```

married
      occupation relationship   race  gender  capital-gain
capital-loss \
0  Machine-op-inspct    Own-child  Black   Male         0
0
1    Farming-fishing     Husband  White   Male         0
0
2    Protective-serv     Husband  White   Male         0
0
3  Machine-op-inspct     Husband  Black   Male       7688
0
4                ?    Own-child  White  Female         0
0

      hours-per-week native-country income
0                40   United-States <=50K
1                50   United-States <=50K
2                40   United-States >50K
3                40   United-States >50K
4                30   United-States <=50K

df = df.replace('?', pd.NA)
df = df.dropna()


```

Distribution of Target Variable

```

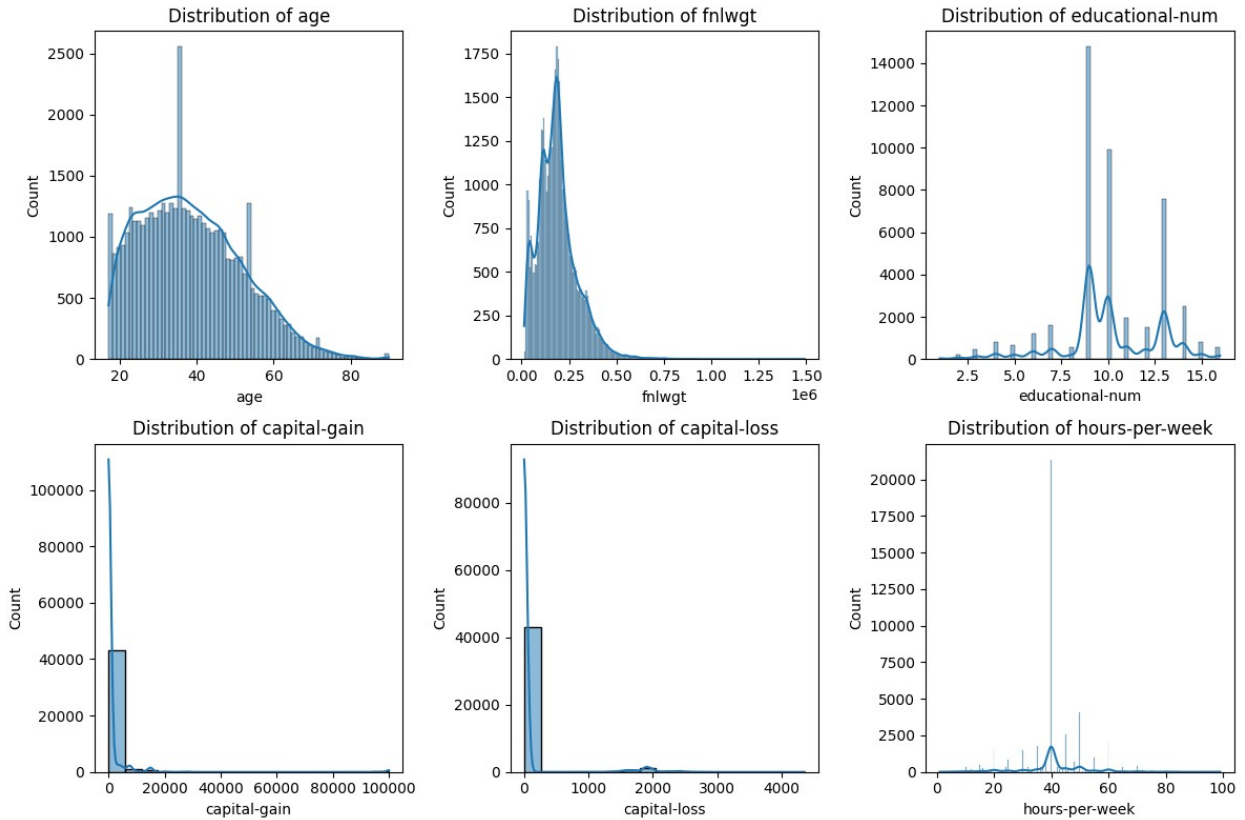
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='income')
plt.title("Distribution of Income")
plt.show()

```



Distribution of Numeric Features

```
numeric_features = ['age', 'fnlwgt', 'educational-num', 'capital-  
gain', 'capital-loss', 'hours-per-week']  
  
plt.figure(figsize=(12, 8))  
for i, feature in enumerate(numeric_features, 1):  
    plt.subplot(2, 3, i)  
    sns.histplot(df[feature], kde=True)  
    plt.title(f'Distribution of {feature}')  
plt.tight_layout()  
plt.show()
```

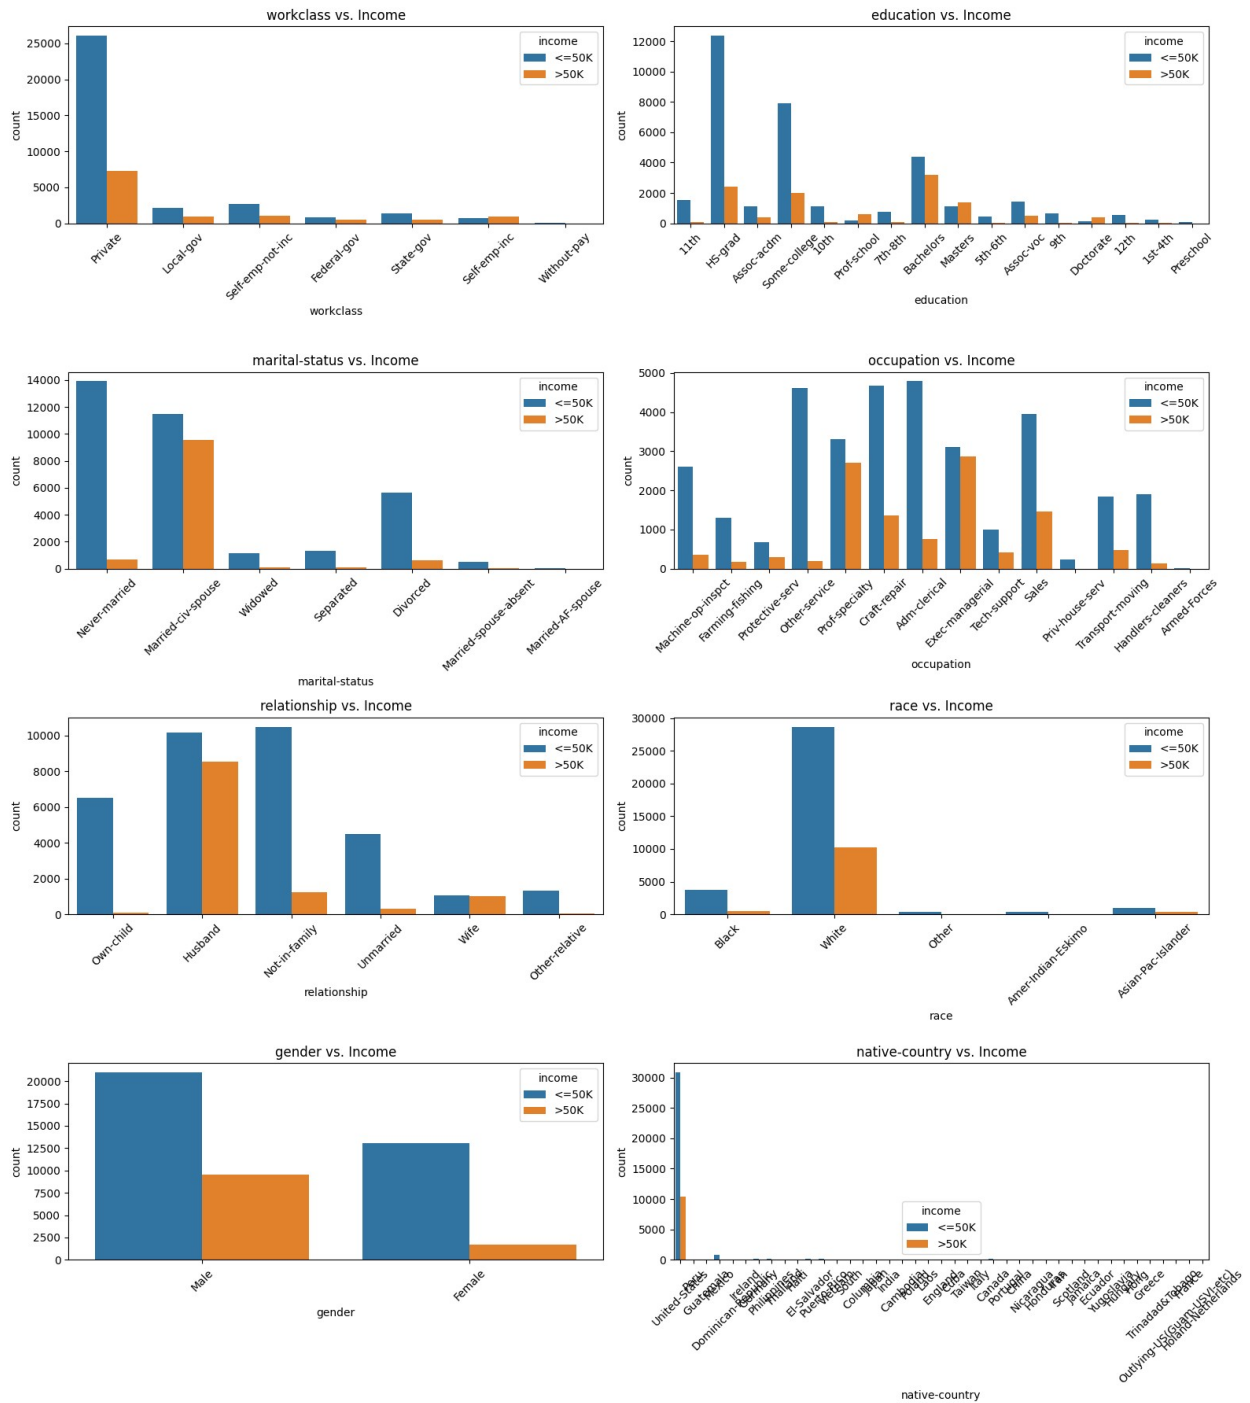


```

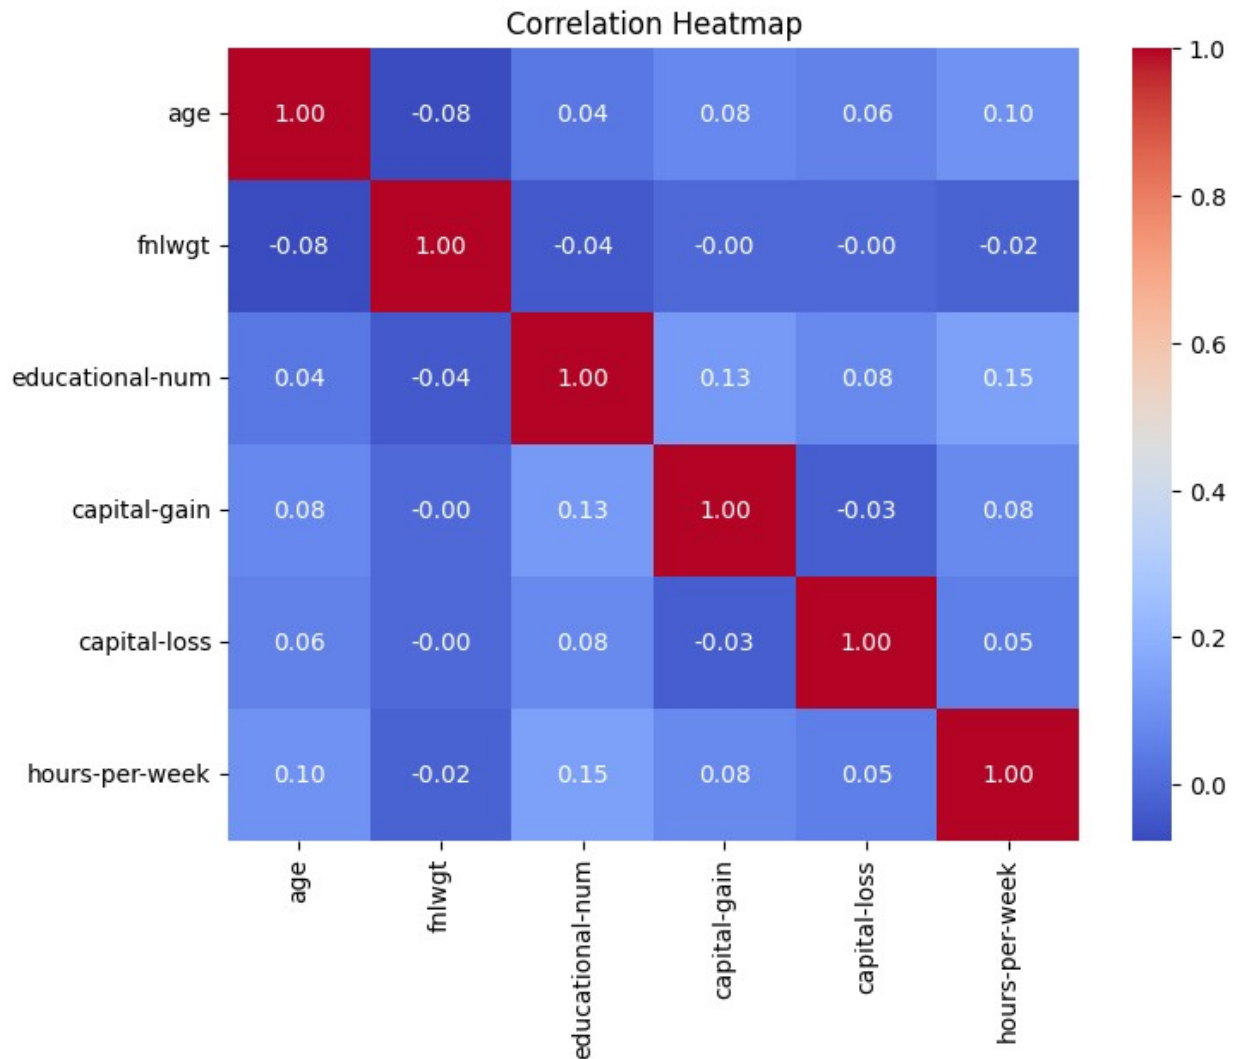
categorical_features = ['workclass', 'education', 'marital-status',
                        'occupation', 'relationship', 'race', 'gender', 'native-country']

plt.figure(figsize=(16, 18))
for i, feature in enumerate(categorical_features, 1):
    plt.subplot(4, 2, i)
    sns.countplot(data=df, x=feature, hue='income')
    plt.title(f'{feature} vs. Income')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```
correlation_matrix = df[numeric_features].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



Feature Engineering

#Example 1: Binning Age

```
age_bins = [0, 25, 45, 65, 100]
age_labels = ['Young', 'Middle-Aged', 'Senior', 'Elderly']
df['Age_Category'] = pd.cut(df['age'], bins=age_bins,
labels=age_labels)
```

Example 2: Education Level Mapping

```
education_mapping = {
    'Preschool': 1, '1st-4th': 2, '5th-6th': 3, '7th-8th': 4, '9th':
5,
    '10th': 6, '11th': 7, '12th': 8, 'HS-grad': 9, 'Some-college': 10,
    'Assoc-acdm': 11, 'Assoc-voc': 12, 'Bachelors': 13, 'Masters': 14,
    'Doctorate': 15}
```

```

}
df['Education_Level'] = df['education'].map(education_mapping)

# Example 3: One-Hot Encoding for Categorical Features

df = pd.get_dummies(df, columns=categorical_features, drop_first=True)

# Example 4: Scaling Numeric Features

scaler = StandardScaler()
numeric_features += ['Education_Level']
df[numeric_features] = scaler.fit_transform(df[numeric_features])

# Example 5: Encoding Income as 0 and 1

df['income'] = df['income'].map({'<=50K': 0, '>50K': 1})
df.head()

```

	age	fnlwgt	educational-num	capital-gain	capital-loss	\
0	-1.024983	0.350889	-1.221559	-0.146733	-0.21878	
1	-0.041455	-0.945878	-0.438122	-0.146733	-0.21878	
2	-0.798015	1.393592	0.737034	-0.146733	-0.21878	
3	0.412481	-0.278420	-0.046403	0.877467	-0.21878	
5	-0.344079	0.084802	-1.613277	-0.146733	-0.21878	

	hours-per-week	income	Age_Category	Education_Level
workclass_Local-gov	\			
0	-0.078120	0	Young	-1.226897
0				
1	0.754701	0	Middle-Aged	-0.417105
0				
2	-0.078120	1	Middle-Aged	0.392686
1				
3	-0.078120	1	Middle-Aged	-0.012210
0				
5	-0.910942	0	Middle-Aged	-1.631793
0				

	...	native-country_Portugal	native-country_Puerto-Rico	\
0	...	0	0	
1	...	0	0	
2	...	0	0	
3	...	0	0	
5	...	0	0	

	native-country_Scotland	native-country_South	native-country_Taiwan	\
0		0	0	
0				
1		0	0	
0				

```

2          0          0
0
3          0          0
0
5          0          0
0

native-country_Thailand native-country_Trinidad&Tobago \
0          0          0
1          0          0
2          0          0
3          0          0
5          0          0

native-country_United-States native-country_Vietnam \
0          1          0
1          1          0
2          1          0
3          1          0
5          1          0

native-country_Yugoslavia
0          0
1          0
2          0
3          0
5          0

[5 rows x 99 columns]

```

Variable Transformations

```

numeric_feature = 'capital-gain'
plt.figure(figsize=(12, 24))

<Figure size 1200x2400 with 0 Axes>


<Figure size 1200x2400 with 0 Axes>

plt.subplot(4, 2, 1)
sns.histplot(df[numeric_feature], kde=True)
plt.title(f'Original {numeric_feature}')

plt.subplot(4, 2, 2)
stats.probplot(df[numeric_feature], dist="norm", plot=plt)
plt.title(f'Q-Q Plot for Original {numeric_feature}')


Text(0.5, 1.0, 'Q-Q Plot for Original capital-gain')

```

```
log_transformed = np.log(df[numeric_feature] + 1)
plt.subplot(4, 2, 3)
sns.histplot(log_transformed, kde=True)
plt.title(f'Logarithm Transformed {numeric_feature}')

plt.subplot(4, 2, 4)
stats.probplot(log_transformed, dist="norm", plot=plt)
# plt.title(f'Q-Q Plot for Logarithm Transformed {numeric_feature}')
Text(0.5, 1.0, 'Q-Q Plot for Logarithm Transformed capital-gain')
```




```
# Square root transformation

sqrt_transformed = np.sqrt(df[numeric_feature])

c:\Users\UMAP\anaconda3\lib\site-packages\pandas\core\
arraylike.py:397: RuntimeWarning: invalid value encountered in sqrt
  result = getattr(ufunc, method)(*inputs, **kwargs)


plt.subplot(4, 2, 5)
sns.histplot(sqrt_transformed, kde=True)
plt.title(f'Square Root Transformed {numeric_feature}')

plt.subplot(4, 2, 6)
stats.probplot(sqrt_transformed, dist="norm", plot=plt)
plt.title(f'Q-Q Plot for Square Root Transformed {numeric_feature}')
Text(0.5, 1.0, 'Q-Q Plot for Square Root Transformed capital-gain')
```




```
reciprocal_transformed = 1 / (df[numeric_feature] + 1)
plt.subplot(4, 2, 7)
sns.histplot(reciprocal_transformed, kde=True)
plt.title(f'Reciprocal Transformed {numeric_feature}')

plt.subplot(4, 2, 8)
stats.probplot(reciprocal_transformed, dist="norm", plot=plt)
plt.title(f'Q-Q Plot for Reciprocal Transformed {numeric_feature}')
Text(0.5, 1.0, 'Q-Q Plot for Reciprocal Transformed capital-gain')
```



```
exp_transformed = np.exp(df[numeric_feature])
plt.subplot(4, 2, 7) # Use 7 for the histogram
sns.histplot(exp_transformed, kde=True)
plt.title(f'Exponential Transformed {numeric_feature}')

plt.subplot(4, 2, 8) # Use 8 for the Q-Q plot
stats.probplot(exp_transformed, dist="norm", plot=plt)
plt.title(f'Q-Q Plot for Exponential Transformed {numeric_feature}')
Text(0.5, 1.0, 'Q-Q Plot for Exponential Transformed capital-gain')
```

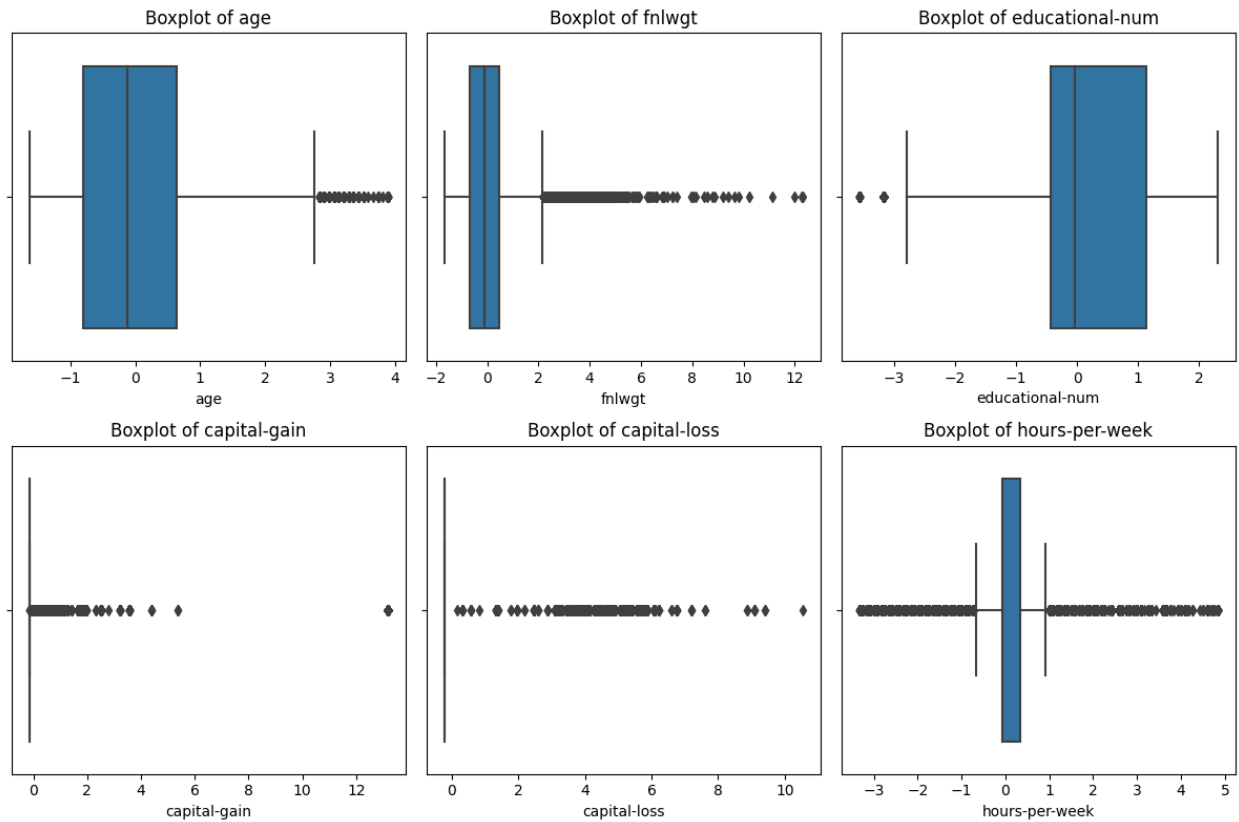


```
X = df.drop('income', axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

numeric_features = ['age', 'fnlwgt', 'educational-num', 'capital-
gain', 'capital-loss', 'hours-per-week']

plt.figure(figsize=(12, 8))
for i, feature in enumerate(numeric_features, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(data=df, x=feature)
    plt.title(f'Boxplot of {feature}')
plt.tight_layout()
plt.show()
```



```
Q1 = df[numeric_feature].quantile(0.25)
Q3 = df[numeric_feature].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Remove outliers

df = df[(df[numeric_feature] >= lower_bound) & (df[numeric_feature] <=
upper_bound)]

# Verify that outliers have been removed

print(f"Number of rows after removing outliers: {len(df)}")

Number of rows after removing outliers: 41432
```